

MINESWEEPER REINFORCEMENT LEARNING WEB APPLICATION

Daniel Short

PROJECT OVERVIEW

- Components
 - Custom Minesweeper Environment
 - Reinforcement Learning (Deep Q-Network) Models
 - Web-based Interactive Application
- Tech Stack
 - JavaScript
 - CSS
 - Python
 - PyTorch
 - Gym
 - Flask
 - Visual Studio

WHAT IS MINESWEEPER?

- **Objective:** Click on all non-mine cells without clicking on a mine.
- **Starting the Game:** Click any cell to begin (first click is safe).
- **Numbers:** Numbers indicate how many mines touch that cell (including diagonally). No number means it isn't touching any mines.
- **Flagging Mines:** Right-click cells that have mines to mark them (*optional*).
- **Using Logic:** Use numbers and flagged cells to logically deduce safe cells.
- **Winning the Game:** Click on all safe cells.
- **Avoiding Mistakes:** Clicking on a mine ends the game immediately.

METHODOLOGY

- Environment & Data
 - Dynamically generated Minesweeper boards
 - First cell click clears a 3x3 area
 - Between 10% and 20% of cells are mines
 - Grid sizes from 5x5 to 10x10
 - Good Example: <https://minesweeperonline.com/>
 - Bad Example: <https://minesweeper.online/game/4524282468>

METHODOLOGY

- Reinforcement Learning Approach
 - Deep Q-Networks (DQN)
 - CNN for Board States
 - Replay buffer
 - Epsilon-greedy policy
 - Curriculum learning strategy

TRAINING AND RESULTS

- Training Setup
 - Network
 - Deep Q-Network (DQN)
 - Double DQN
 - Dueling DQN
 - Board State
 - Max Pooling
 - Adaptive Pooling
 - Global Pooling
 - Memory Replay Buffer
 - Standard Replay Buffer
 - Prioritized Replay Buffer

TRAINING AND RESULTS

- Hyperparameters
 - Episodes = 10,000
 - Batch size = 64
 - Train Frequency = 50 episodes
 - Update Target Model Frequency = 1,000 episodes
 - Success Threshold = 50%
- Performance Metrics
 - Success Rate

DQN MODEL POINT VALUES

- Positive Rewards
 - Open a cell correctly: +0.3 points
 - Forced guess (correct): +0.3 points
 - Forced guess (incorrect): +0.3 points
 - Win the game: +1.0 points
- Negative Rewards
 - Guess (correct): -0.3 points
 - Guess (incorrect): -0.3 points
 - Lose the game: -1.0 points

DEMO OF THE WEB APPLICATION

- Interactive Web Interface
 - AI Gameplay
 - Visual gameplay of the RL agent
 - User Gameplay
 - User-selectable grid sizes

CHALLENGES AND SOLUTIONS

- Challenges
 - Balancing exploration vs. exploitation
 - Managing computation complexity
 - Finding point values to incentivize learning
 - Determining what is determined successful performance in training
- Solutions
 - Improved epsilon decay strategy
 - Curriculum learning for stability
 - Trial and error

LEARNING OUTCOMES

- Learning Outcomes
 - Experience with reinforcement learning (big topic with ChatGPT recently)
 - Custom environment creation
 - Model Deployment and UI designing

CONCLUSION AND FUTURE WORK

- Project Results
 - Double DQN outperformed
 - Adaptive Pooling outperformed
 - Standard Replay Buffer outperformed
- Project Achievements
 - Successful RL agent training
 - Effective curriculum learning implementation
 - Interactive and user-friendly web application
- Future Directions
 - Implement Visual Transformers instead of CNNs
 - Expand to similar games like Nonogram
 - Optimize model parameters